# WHEN TASK-SPECIFIC LEARNING OUTPERFORMS TRANSFER LEARNING: A BENCHMARK OF GENE AND EXPRESSION ENCODING STRATEGIES*

**Igor Sadalski**
Somite.ai
Boston, MA, USA
igor.sadalski@gmail.com

## ABSTRACT

Single-cell foundational models use transformer backbones that require careful engineering of gene and expression encoding strategies, yet there is no consensus on which encoding techniques are effective. We systematically compare encoding paradigms by training models from scratch under controlled conditions, evaluating them using a comprehensive evaluation suite spanning multiple evaluation dimensions: batch correction metrics (integration local inverse Simpson's index, graph connectivity, average silhouette width), biological preservation metrics (normalized mutual information, adjusted rand index, isolated label metrics), classification performance (accuracy and macro F1), and reconstruction quality. We scale pretraining to 10 million cells across 100 diverse datasets (tenfold increase over similar studies) and compare learned versus pretrained protein model gene embeddings with four expression encoding methods (discrete binning, differentiable soft binning, logarithmic binning, and continuous encoding). Contrary to common assumptions, task-specific learned embeddings consistently outperform pretrained protein model embeddings across all metrics. The best configuration (differentiable soft binning with learned embeddings) achieves superior classification performance and substantial improvements on batch correction and biological preservation metrics. Our work provides clear empirical guidance for model design and establishes a systematic benchmark for encoding strategies in single-cell foundational models.

## 1 INTRODUCTION

A major goal in recent years has been to build an AI virtual cell, i.e., multi-scale, multi-modal neural network models that can represent and simulate cellular behaviour across diverse states (Bunne et al., 2024). Leading models promise to enable universal embeddings (Rosen et al., 2023), cross-species transfer (Pearce et al., 2025), multi-task transfer learning for downstream applications (Cui et al., 2024; Theodoris et al., 2023), and batch correction (Wang et al., 2021). While foundational models can be generated for different omics data types (e.g., protein, transcriptomics, tissues), here we focus on transcriptomic data. In this domain, most of the models (e.g. Cui et al. (2024); Adduri et al. (2025); Pearce et al. (2025)) use a transformer (Vaswani et al., 2017) backbone. Since transcriptomic data inherently consists of two pieces of information, gene identities and their expressions, architects of these models are faced with a design choice on how to encode this information into embeddings that transformers can process.

Different encoding strategies embody distinct assumptions about what information is most useful. For genes, it's task-specific patterns versus biological priors; for expression values, it's continuous precision versus discrete robustness. Figure 4 provides a schematic overview of representative encoding paradigms. For gene encoding, some models learn embeddings de novo to capture dataset-specific co-expression relationships (Cui et al., 2024), while others leverage transfer learning from pretrained protein language models such as ESM-2 (Lin et al., 2023; Adduri et al., 2025) to incorpo-

*draft in preparation for ICLR 2026, 2nd Workshop on Foundation Models for Science

rate evolutionary biological knowledge. For expression encoding, strategies range from discretizing into bins for computational efficiency (Cui et al., 2024; Gandhi et al., 2025) to using continuous or soft binning via MLPs (Pearce et al., 2025; Ho et al., 2024; Adduri et al., 2025) or applying binning based on logarithmic transformation (Wang et al., 2021).

We present a controlled benchmark that quantifies these methodologies using consistent model architecture and training procedures at a large scale. We overcome the limitations of previous work by scaling up the pretraining data tenfold to 10 million cells, increasing dataset diversity by using 100 different datasets, introducing new tokenization strategies such as log binning and raw embedding, and performing evaluation against the Tabula Sapiens v2 benchmark.

The main contributions of this work are:

- **A comprehensive evaluation suite** spanning batch correction, biological preservation, classification, and reconstruction metrics across 26 diverse tissue datasets.
- **A tenfold increase in scale of pretraining data** to 10 million cells across 100 diverse datasets.
- **A systematic evaluation of encoding strategies** comparing learned versus pretrained protein model gene embeddings and four expression encoding methods (discrete binning, differentiable soft binning, logarithmic binning, and continuous encoding) under controlled conditions with consistent model architecture and training procedures.

## 2 Related Work

Limited scientific literature has addressed this problem. Most benchmarking efforts have focused on evaluating downstream applications, such as perturbation prediction (Ahlmann-Eltze et al., 2024; Wenteler et al., 2024), using already pretrained models. In this work, we focus more on specific architectural choices and pretrain our models from scratch. A similar approach was taken in HEIMDALL (Haber et al., 2025), a modular tokenization framework, where authors tried to evaluate different encoding strategies. However, they pretrained their transformer model on only a few datasets with a total of 1 million cells. This limited scale and dataset diversity likely prevents transformers from learning robust, generalizable representations, as transformer performance scales with training data size (Kaplan et al., 2020). In comparison, regular transformer models in the field are trained on, e.g., 266 million cells (Gandhi et al., 2025) and hundreds of diverse datasets. Given improved transformer performance with data scale (Kaplan et al., 2020), we scale our pretraining to better align with the field's standard practices.

## 3 Methods

### 3.1 Model Architecture and Pretraining

We represented cells as bags of gene-expression pairs. Each cell $i$ contains $M_i$ genes with non-zero expression values, represented as a sequence of gene-expression pairs:

$$\mathcal{C}_i = ((g_{i,1}, x_{i,1}), \ldots, (g_{i,M_i}, x_{i,M_i})) \tag{1}$$

where $g_{i,j}$ denotes the $j$-th gene identifier and $x_{i,j}$ denotes its corresponding expression value in cell $i$.

Prior to encoding, we normalized expression values for each cell to a constant total $c_{\text{norm}}$ and applied log1p transformation:

$$\tilde{x}_{i,j} = \log\left(1 + \frac{x_{i,j}}{\sum_{k=1}^{M_i} x_{i,k}} \times c_{\text{norm}}\right) \tag{2}$$

where $M_i$ is the number of genes in cell $i$, $c_{\text{norm}}$ is a normalization constant (typically $10^4$ or $10^6$), and $k$ indexes over all genes in cell $i$. This normalization step ensures consistent scaling across cells with varying sequencing depths.

Because transformers scale with the square of the size of the context window (number of inputs), we employed a sampling strategy to select a subset of genes for each cell. Let $S : \mathcal{C}_i \rightarrow$

$\{(g_{i,j}, x_{i,j})\}_{j=1}^{K}$ be a sampling function that takes cell $i$ and returns $K$ gene-expression pairs, where $K = \min(\text{context\_window}, M_i)$. During training, for datasets where the number of non-zero expressed genes exceeded the context window, we randomly sampled $K$ genes from all non-zero expressed genes in each cell:

$$\mathcal{S}^{(i)}_{\text{random}} = \text{RandomSample}(\{g_j : x_{i,j} > 0\}, K) \tag{3}$$

This approach ensures diverse gene representation across training examples while maintaining computational feasibility for large gene vocabularies.

The training objective was to predict masked expression values given the gene identities and unmasked expression context. To this end, we first computed embeddings for both gene identities and expression values. The gene embedding for each gene was obtained as:

$$\boldsymbol{e}_g^{(i,j)} = \text{enc}_g(g_{i,j}) \in \mathbb{R}^{d_g}, \tag{4}$$

where $\text{enc}_g$ is the gene encoding function (detailed in Section 3.2), $d_g$ is the gene embedding dimension, and $\boldsymbol{e}_g^{(i,j)} \in \mathbb{R}^{d_g}$ is the gene embedding vector. For expression values, we randomly masked a fraction $p_{\text{mask}}$ of gene-expression pairs during training. The expression embedding was computed as:

$$\boldsymbol{e}_x^{(i,j)} = \begin{cases} \text{enc}_x(\tilde{x}_{i,j}) & \text{with probability } 1 - p_{\text{mask}} \\ \boldsymbol{m} & \text{with probability } p_{\text{mask}} \end{cases} \tag{5}$$

where $\text{enc}_x$ is the expression encoding function (detailed in Section 3.3), $\boldsymbol{e}_x^{(i,j)} \in \mathbb{R}^{d_x}$ is the expression embedding vector, $p_{\text{mask}}$ is the masking probability, and $\boldsymbol{m} \in \mathbb{R}^{d_x}$ is a learnable mask token with dimension $d_x$ matching the expression embedding dimension.

The combined gene-expression embedding for the $j$-th gene in cell $i$ is obtained by summing the gene and expression embeddings:

$$\boldsymbol{z}_0^{(i,j)} = \boldsymbol{e}_g^{(i,j)} + \boldsymbol{e}_x^{(i,j)} \tag{6}$$

where $\boldsymbol{z}_0^{(i,j)} \in \mathbb{R}^d$ is the combined embedding for gene $j$ in cell $i$, and $d$ is the model dimension (equal to both $d_g$ and $d_x$). The combined embeddings for all $K$ genes in cell $i$ are concatenated together to form the input sequence for the model:

$$\boldsymbol{z}_0^{(i)} = [\boldsymbol{z}_0^{(i,1)}, \ldots, \boldsymbol{z}_0^{(i,K)}] \tag{7}$$

where $\boldsymbol{z}_0^{(i)} \in \mathbb{R}^{K \times d}$ is the input sequence for cell $i$ containing $K$ gene embeddings. These combined embeddings were then passed through a transformer encoder. Grouping the gene-embeddings per cell and recursively applying the transformer layers, we obtained:

$$\boldsymbol{z}_l^{(i)} = f_{\text{transformer}}(\boldsymbol{z}_{l-1}^{(i)}) \tag{8}$$

where $f_{\text{transformer}}$ denotes a standard transformer encoder layer, $l \in \{1, \ldots, n\}$ indexes the layers, and $n$ is the total number of transformer layers.

The predicted expression value for the $j$-th gene in cell $i$ is decoded from the final transformer layer output using a simple MLP:

$$\bar{x}_{i,j} = \text{MLP}(\boldsymbol{z}_n^{(i,j)}) \tag{9}$$

where $\boldsymbol{z}_n^{(i,j)} \in \mathbb{R}^d$ is the output embedding for gene $j$ in cell $i$ from the final transformer layer, and $\bar{x}_{i,j} \in \mathbb{R}$ is the predicted expression value. We optimized the model using mean squared error loss for reconstruction of masked gene expression values (Wang et al., 2021; Adduri et al., 2025; Ho et al., 2024; Cui et al., 2024):

$$\mathcal{L}_{i,j} = \frac{1}{|\mathcal{U}_{\text{unk}}|} \sum_{j \in \mathcal{U}_{\text{unk}}} (\tilde{x}_{i,j} - \bar{x}_{i,j})^2 \tag{10}$$

where $\mathcal{L}_{i,j}$ is the loss for gene $j$ in cell $i$, $\mathcal{U}_{\text{unk}}$ represents the set of masked gene indices, $\tilde{x}_{i,j}$ is the normalized true expression value, and $\bar{x}_{i,j}$ is the predicted expression value from Equation 9.

### 3.2 GENE ENCODING STRATEGIES

**Learned encoding** is obtained by passing the gene identifier through a learned embedding table (Cui et al., 2024; Gandhi et al., 2025). This approach allows the model to generate task-specific representations for each gene, enabling it to capture dataset-specific relationships.

$$\text{enc}_g^{\text{learned}}(g_{i,j}) = \text{Embedding}(g_{i,j}) \tag{11}$$

**ESM-2 encoding** uses precomputed embeddings retrieved from the ESM-2 (3B) model dictionary (Lin et al., 2023). STATE (Adduri et al., 2025) leverages this approach, using large pretrained protein language models to provide biologically-informed representations that are consistent across datasets and species, making it a strong choice for transfer learning and handling new or rare gene symbols. We projected these embeddings through an MLP to match the model dimension.

$$\text{enc}_g^{\text{ESM-2}}(g_{i,j}) = \text{MLP}(\mathcal{E}(g_{i,j})) \tag{12}$$

where $\mathcal{E}(g_{i,j})$ denotes the precomputed ESM-2 embedding for gene identifier $g_{i,j}$, and MLP is a multi-layer perceptron that projects the embedding to the model dimension.

### 3.3 EXPRESSION ENCODING STRATEGIES

**Raw expression encoding** uses a simple MLP on normalized data. This approach preserves the full, continuous information from the expression measurement and is the most direct way to encode quantitative gene expression levels.

$$\text{enc}_x^{\text{raw}}(\tilde{x}_{i,j}) = \text{MLP}(\tilde{x}_{i,j}) \tag{13}$$

**Hard binning encoding** discretizes expression values into bins. This method groups expression levels into discrete intervals, trading off resolution for robustness and simplifying the input space (Cui et al., 2024; Gandhi et al., 2025).

$$b_{i,j} = \begin{cases} k, & \text{if } x_{i,j} > 0 \text{ and } x_{i,j} \in [\beta_k, \beta_{k+1}], \\ 0, & \text{if } x_{i,j} = 0, \end{cases} \tag{14}$$

where $b_{i,j}$ is the bin index for the expression value of gene $j$ in cell $i$, $k$ is the bin index, and $\beta_k$ and $\beta_{k+1}$ are the lower and upper boundaries of bin $k$, respectively. Expression embedding is obtained using an embedding layer:

$$\text{enc}_x^{\text{hard bin}}(x_{i,j}) = \text{Embedding}(b_{i,j}) \tag{15}$$

**Log binning encoding** compresses a wide dynamic range of expression values using a logarithmic transformation before discretizing. This approach can potentially mitigate the effects of outliers or skewed distributions. STATE (Adduri et al., 2025) uses this approach with ESM-2 embeddings, while scBERT (Wang et al., 2021) employs log binning with discrete tokenization.

$$b_{i,j} = \min\left(\left\lfloor \log_2\left(x_{i,j} + 1\right)\right\rfloor, B_{\max}\right) \tag{16}$$

where $b_{i,j}$ is the bin index for the expression value of gene $j$ in cell $i$, and $B_{\max}$ is the maximum bin index. Bins are embedded using an embedding layer:

$$\text{enc}_x^{\text{log bin}}(x_{i,j}) = \text{Embedding}(b_{i,j}) \tag{17}$$

**Soft binning encoding** uses a softmax over potential bins to allow fractional/bin-weighted expression, which can capture uncertainty and subtle intensity differences between expression values, making the encoding differentiable and potentially more expressive (Ho et al., 2024; Hao et al., 2024; Pearce et al., 2025).

$$\boldsymbol{\alpha}_x^{(i,j)} = \text{Softmax}\left(\boldsymbol{W}_{x,2}\text{LeakyReLU}\left(\boldsymbol{W}_{x,1}\, x_{i,j}\right)\right) \tag{18}$$

where $\boldsymbol{\alpha}_x^{(i,j)}$ is a vector of bin weights for the expression value of gene $j$ in cell $i$, $\boldsymbol{W}_{x,1}$ and $\boldsymbol{W}_{x,2}$ are learnable weight matrices, and LeakyReLU is the LeakyReLU activation function. The expression embedding is obtained by performing a soft lookup in the embedding table:

$$\text{enc}_x^{\text{soft bin}}(x_{i,j}) = \sum_{k=1}^{b} \alpha_{x,k}^{(i,j)} \boldsymbol{T}_k \tag{19}$$

where $b$ is the number of bins, $\alpha_{x,k}^{(i,j)}$ is the $k$-th element of $\boldsymbol{\alpha}_x^{(i,j)}$, and $\boldsymbol{T}_k$ is the $k$-th embedding vector in the embedding table $\boldsymbol{T}$.

# 4 EXPERIMENTS

## 4.1 TRAINING

We trained our models on 101 diverse single-cell RNA sequencing datasets comprising over 10 million cells, randomly selected from the curated collection used to train the Transcriptformer model (Pearce et al., 2025). The datasets span diverse tissue types, experimental protocols, species, biological conditions, developmental stages, and disease states (see Appendix A). We evaluated four different expression encoding strategies and two different gene encoding strategies by training models for all combinations of these strategies (8 models total). During the training process we used the hyperparameters defined in Table 1 in the Appendix, which were based on previous works (Cui et al., 2024; Adduri et al., 2025; Pearce et al., 2025). All models were trained using automatic mixed precision (AMP) with FP16 to accelerate training. During training, we recorded the reconstruction loss for each model using Mean Squared Error on masked gene expression values, which directly measures the model's ability to predict masked expression values during pretraining. We recorded the best model checkpoints based on the validation reconstruction loss.

## 4.2 EVALUATION

We evaluated all models on 26 tissue-specific datasets from the Tabula Sapiens v2 benchmark (Tabula Sapiens Consortium et al., 2022), comprising more than half a million cells across diverse human tissues. This benchmark represents a standard evaluation protocol in the field and has been used to evaluate other foundational models (Pearce et al., 2025). Using each trained model, we embedded 26 datasets that constituted the evaluation dataset. For each embedded evaluation dataset, we generated cell embeddings by extracting token embeddings from the final encoder layer (where each token represents a gene-expression pair) and mean-pooling these embeddings across all valid (non-padded) positions for each cell. Our embeddings were annotated with cell type. Additionally, for the batch key we used the `10X_run` identifier, which distinguishes cells from different 10X Genomics sequencing runs.

**Batch correction metrics.** To assess batch effect removal, we used the scib package (Luecken et al., 2022) and computed: Integration Local Inverse Simpson's Index, which measures the diversity of batch labels in the local neighborhood of each cell to quantify batch mixing; Cell-type Local Inverse Simpson's Index, which evaluates the preservation of biological structure by measuring the diversity of cell-type labels in the local neighborhood; Average Silhouette Width for batch, which measures how well batches are mixed globally using average silhouette width; and graph connectivity, which assesses whether cells sharing the same label form a fully connected subgraph in the k-nearest neighbors graph.

**Biological preservation metrics.** Similarly to batch correction metrics (also using the scib package (Luecken et al., 2022)), we computed: Average Silhouette Width for label, which quantifies how well cells of the same cell type cluster together using average silhouette width computed over cell-type labels; Normalized Mutual Information, which measures the normalized mutual information between predicted clusters and true cell-type labels; Adjusted Rand Index, which computes the adjusted rand index between predicted clusters and true cell-type labels; isolated label F1 score, which evaluates the F1 score for cell types that are isolated in the embedding space; and isolated label silhouette, which measures the silhouette score for isolated cell types.

**Cell type classification.** To compute cell type classification performance, we extracted cell embeddings by mean-pooling token embeddings from the final encoder layer (where each token represents a gene-expression pair) (Pearce et al., 2025). We filtered out cells with missing cell type annotations and removed cell types with fewer than 250 cells to ensure robust evaluation. We split the data into training and test sets using an 80/20 stratified split. We trained a k-nearest neighbors classifier with $k = 10$ on the training set and evaluated performance on the held-out test set.

# 5 RESULTS

We evaluated all 8 model configurations (2 gene encoding strategies $\times$ 4 expression encoding strategies) on 26 tissue-specific datasets from the Tabula Sapiens v2 benchmark. For each encoder config-

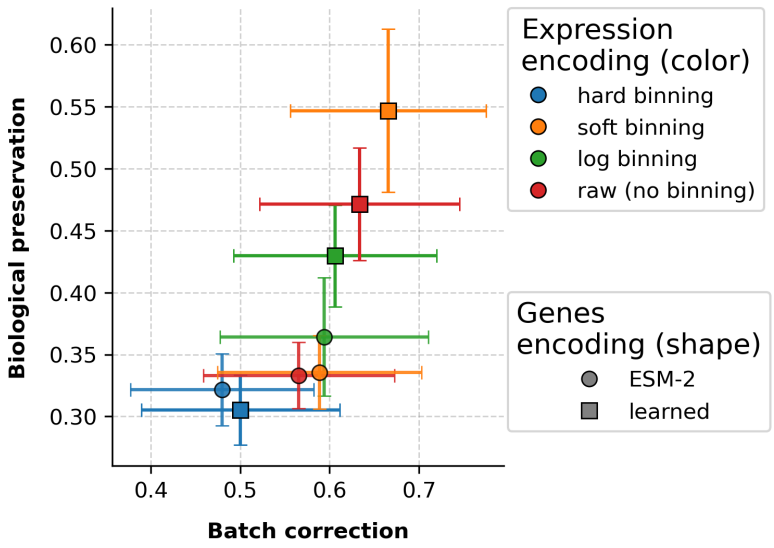


Figure 1: Performance comparison across encoding strategies on batch correction and biological preservation metrics. To obtain a single point, first, we take one of the evaluation datasets (we subset 11 out of 26 datasets). For that dataset we compute several evaluation metrics, then averaged those metrics within the dataset to obtain a per-dataset mean for each model configuration. Next, we aggregated results by taking the mean and standard deviation of these per-dataset metric means across all datasets (so error bars are standard deviations of the per-dataset means). The X-axis shows batch removal performance, which is the average of metrics related to batch correction (Integration Local Inverse Simpson's Index, graph connectivity, Average Silhouette Width for batch). The Y-axis shows biology conservation performance, which is the average of metrics that capture biological signal preservation (Average Silhouette Width for label, Cell-type Local Inverse Simpson's Index, Normalized Mutual Information, Adjusted Rand Index, isolated label F1 score, isolated label silhouette score). Colors distinguish expression encoding strategies (hard binning, soft binning, log binning, raw expression), and marker shapes indicate gene encoding strategies (learned embedding or ESM-2 embedding).

uration, we embedded separately each of the 26 datasets and then computed selected metrics from the scib package (Luecken et al., 2022) which are used to evaluate batch correction and biological preservation, cell type classification performance, and pretraining reconstruction loss. Our evaluation revealed three key findings: (1) learned gene embeddings substantially outperformed ESM-2 embeddings across all metrics, (2) soft binning emerged as the optimal expression encoding strategy, and (3) pretraining reconstruction quality correlated with downstream task performance.

### 5.1 Batch Correction and Biological Preservation

Figure 1 shows a comparison of the trained models' ability to perform batch correction and preserve biological features. Each point summarizes the model's average batch correction (average of selected metrics) and biological preservation (average of selected metrics) performance aggregated across datasets. In these results, learned gene embeddings with soft binning achieve the best performance, representing a 16% improvement over the second-best configuration (raw encoding with learned embeddings at 0.471). Overall, learned gene embeddings provide a 29% relative improvement over ESM-2 embeddings on the combined metrics, supporting the advantage of using task-specific learned representations for both batch correction and biological signal conservation. The batch correction metrics show high variability, with an average coefficient of variation of 29% (mean standard deviation 0.11 over mean score 0.38), reflecting substantial differences across the 26 evaluation datasets.

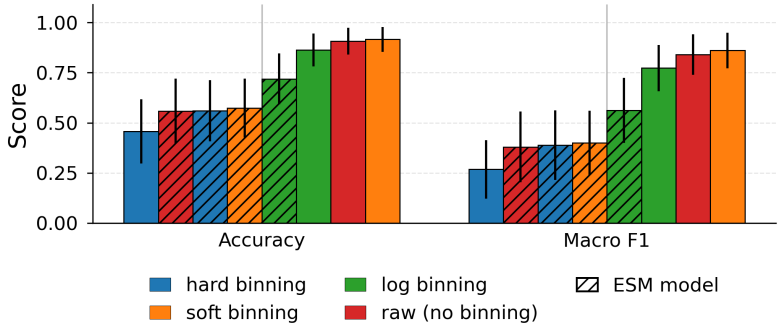


Figure 2: Cell type classification performance across encoding strategies. Bars show mean accuracy and macro F1 score with error bars indicating standard deviation across 26 evaluation datasets. Colors indicate expression encoding strategies (hard, soft, log, raw); hatching (///) indicates ESM-2 gene encoding. Configurations are sorted in ascending order.

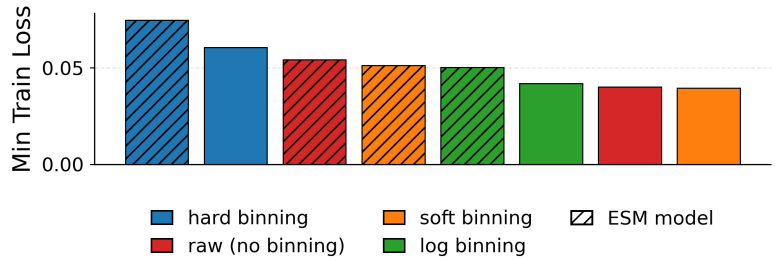


Figure 3: Minimum Mean Squared Error (MSE) training loss achieved during pretraining for each model configuration. Bars are sorted by minimum loss in descending order. Colors indicate expression encoding strategies (hard, soft, log, raw); hatching (///) indicates ESM-2 gene encoding.

### 5.2 Cell Type Classification Performance

Figure 2 shows cell type classification performance across all encoding configurations. Learned gene embeddings outperform ESM-2 embeddings for three out of four cases, achieving 30% higher accuracy and 58% higher macro F1 score on average. The best-performing configuration combines learned gene embeddings with soft binning, achieving 91.6% accuracy and 86.0% macro F1 score. This represents a 92% improvement in macro F1 over hard binning with learned embeddings (44.8% macro F1). The consistent advantage of learned embeddings across all expression encoding strategies suggests that task-specific representations capture more relevant information for transcriptomic tasks than transfer learning from protein language models. Furthermore, the order of performance for different methods, in both accuracy and macro F1, is mostly consistent with other results **??**.

### 5.3 Pretraining Reconstruction Quality

Figure 3 shows the minimum training reconstruction loss achieved during pretraining for each model configuration. Configurations with learned gene embeddings achieve lower minimum training loss than ESM-2 configurations (0.046 vs 0.058 MSE on average, 21% improvement), aligning with downstream performance. This correlation suggests that better pretraining reconstruction quality translates to improved downstream task performance. The higher training loss for ESM-2 configurations, combined with observed training instability (loss diverging in all four ESM-2 cases and learned hard binning), indicates that transfer learning from protein language models may introduce optimization challenges that limit effectiveness for transcriptomic tasks.

## 6 DISCUSSION

**Learned gene embeddings outperform ESM-2 encoding.** Despite an MLP layer projecting ESM-2 embeddings to the model dimension, models trained with ESM-2 embeddings showed unstable training with loss diverging in all four cases. Learned embeddings achieved an average 39% relative improvement over ESM-2 embeddings across all metrics. This gap suggests that task-specific representations captured dataset-specific co-expression relationships more relevant for transcriptomic tasks than evolutionary protein-level information. Additionally, single-cell data includes many non-coding genes, pseudogenes, and gene isoforms not well-represented in protein language models.

**Soft binning provides optimal expression encoding.** Soft binning balances discrete and continuous representations. When combined with learned gene embeddings, it achieved 91.6% classification accuracy and 86.0% macro F1 score, with a 41% improvement on batch integration and biological preservation metrics and 92% improvement on macro F1 over hard binning. Unlike hard binning, soft binning preserved information at bin boundaries and captured subtle expression differences, while providing regularization that may help learn more robust representations than raw continuous encoding.

Our evaluation is limited to human tissues, cell type classification, and a single model architecture. Future work should evaluate additional downstream tasks and species, investigate alternative architectures, and provide mechanistic insights into why learned embeddings and soft binning outperform their alternatives.

## 7 CONCLUSION

We present a large-scale systematic benchmark comparing gene and expression encoding strategies, training 8 configurations on over 10 million cells across 100 diverse datasets under controlled conditions. We evaluate all configurations using a comprehensive evaluation suite spanning batch correction, biological preservation, classification, and reconstruction metrics across 26 diverse tissue datasets. Our evaluation yields three key findings: (1) task-specific learned gene embeddings substantially outperform pretrained protein model embeddings, achieving 39% average improvement across all metrics; (2) differentiable soft binning emerges as the optimal expression encoding strategy when combined with learned embeddings, achieving 91.6% classification accuracy and 86.0% macro F1 score; (3) gene encoding choice has greater impact than expression encoding choice, with learned embeddings providing consistent improvements across all expression encoding strategies. This benchmark establishes clear empirical guidance for model design decisions and provides reproducible results for future method development.

## IMPACT STATEMENT

This paper presents a benchmark study that provides empirical guidance for designing effective foundational models in single-cell RNA sequencing analysis. Potential positive impacts include enabling more effective analysis tools for biomedical research and providing reproducible benchmarks for method comparison. Potential negative impacts are limited, as this is a methodological study focused on model architecture rather than direct applications.

## ACKNOWLEDGEMENTS

We thank Somite AI for their help and guidance throughout this work.

## REFERENCES

Abhinav K Adduri, Dhruv Gautam, Beatrice Bevilacqua, Alishba Imran, Rohan Shah, Mohsen Naghipourfar, Noam Teyssier, Rajesh Ilango, Sanjay Nagaraj, Mingze Dong, Chiara Ricci-Tam, Christopher Carpenter, Vishvak Subramanyam, Aidan Winters, Sravya Tirukkovular, Jeremy Sullivan, Brian S Plosky, Basak Eraslan, Nicholas D Youngblut, Jure Leskovec, Luke A Gilbert, Silvana Konermann, Patrick D Hsu, Alexander Dobin, Dave P Burke, Hani Goodarzi, and Yusuf H

Roohani. Predicting cellular responses to perturbation across diverse contexts with state. *bioRxiv*, 2025. doi: 10.1101/2025.06.26.661135. URL https://doi.org/10.1101/2025.06.26.661135. Preprint.

Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep learning-based predictions of gene perturbation effects do not yet outperform simple linear methods. *Nature Methods*, 2024. doi: 10.1038/s41592-025-02772-6. URL https://doi.org/10.1038/s41592-025-02772-6. Published version.

Charlotte Bunne, Yusuf Roohani, Yanay Rosen, others, Emma Lundberg, Jure Leskovec, and Stephen R Quake. How to build the virtual cell with artificial intelligence: Priorities and opportunities. *Cell*, 187(25), 2024.

Haotian Cui, Chloe Wang, Hassaan Maan, Kuan Pang, Fengning Luo, Nan Duan, and Bo Wang. scgpt: toward building a foundation model for single-cell multi-omics using generative ai. *Nature Methods*, 21:1–12, 2024. doi: 10.1038/s41592-024-02201-0. URL https://doi.org/10.1038/s41592-024-02201-0.

Shreshth Gandhi, Farnoosh Javadi, Valentine Svensson, Umair Khan, Matthew G. Jones, John Yu, Daniele Merico, Hani Goodarzi, and Nima Alidoust. Tahoe-x1: Scaling perturbation-trained single-cell foundation models to 3 billion parameters. *bioRxiv*, 2025. doi: 10.1101/2025.10.23.683759. URL https://doi.org/10.1101/2025.10.23.683759. Preprint.

Ellie Haber, Shahul Alam, Nicholas Ho, Renming Liu, Evan Trop, Shaoheng Liang, Muyu Yang, Spencer Krieger, and Jian Ma. Heimdall: A modular framework for tokenization in single-cell foundation models. *bioRxiv*, pp. 2025–11, 2025.

M. Hao, J. Gong, X. Zeng, et al. Large-scale foundation model on single-cell transcriptomics. *Nature Methods*, 21:1481–1491, 2024. doi: 10.1038/s41592-024-02305-7. URL https://doi.org/10.1038/s41592-024-02305-7.

Nicholas Ho, Caleb N Ellington, Jinyu Hou, Sohan Addagudi, Shentong Mo, Tianhua Tao, Dian Li, Yonghao Zhuang, Hongyi Wang, Xingyi Cheng, Le Song, and Eric P Xing. Scaling dense representations for single cell with transcriptome-scale context. *bioRxiv*, 2024. doi: 10.1101/2024.11.28.625303. URL https://doi.org/10.1101/2024.11.28.625303. Preprint.

Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020. URL https://arxiv.org/abs/2001.08361. Preprint.

Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaron Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of protein structure from sequence. *Nature*, 622(7983):827–835, 2023. doi: 10.1038/s41586-023-06520-w. URL https://doi.org/10.1038/s41586-023-06520-w.

Malte D. Luecken, M. Büttner, K. Chaichoompu, A. Danese, M. Interlandi, M. F. Mueller, D. C. Strobl, L. Zappia, M. Dugas, M. Colomé-Tatché, and Fabian J. Theis. Benchmarking atlas-level data integration in single-cell genomics. *Nature Methods*, 19:41–50, 2022. doi: 10.1038/s41592-021-01336-8. URL https://doi.org/10.1038/s41592-021-01336-8.

James D Pearce, Sara E Simmonds, Gita Mahmoudabadi, Lakshmi Krishnan, Giovanni Palla, Ana-Maria Istrate, Alexander Tarashansky, Benjamin Nelson, Omar Valenzuela, Donghui Li, Stephen R Quake, and Theofanis Karaletsos. A cross-species generative cell atlas across 1.5 billion years of evolution: The transcriptformer single-cell model. *bioRxiv*, 2025. doi: 10.1101/2025.04.25.650731. URL https://doi.org/10.1101/2025.04.25.650731. Preprint.

Yanay Rosen, Yusuf Roohani, Ayush Agarwal, Leon Samotorčan, Tabula Sapiens Consortium, Stephen R Quake, and Jure Leskovec. Universal cell embeddings: A foundation model for cell biology. *bioRxiv*, 2023. doi: 10.1101/2023.11.28.568918. URL https://doi.org/10.1101/2023.11.28.568918. Preprint.

Tabula Sapiens Consortium, Robert C Jones, Jim Karkanias, Mark A Krasnow, Angela Oliveira Pisco, Stephen R Quake, Julia Salzman, Nir Yosef, Bryan Bulthaup, Phillip Brown, Will Harper, Michael Hemenez, Ravikumar Ponnusamy, Ahmad Salehi, Bhavani A Sanagavarapu, Eileen Spallino, Kalleen A Aaron, Waldo Concepcion, James M Gardner, Brendan Kelly, Nicholas Neidlinger, Zifa Wang, Sheela Crasta, Saroja Kolluru, Maurizio Morri, Yan Tan, Kyle J Travaglini, Chenling Xu, Maria Alimova, Nicholas E Banovich, Ben A Barres, Philip A Beachy, Biter Bilen, Douglas Brownfield, Charles K F Chan, Songming Chen, Michael F Clarke, Sabrina D Conley, Spyros Darmanis, Aaron Demers, Kubilay Demir, Antoine de Morree, Tony Divita, Haley du Bois, Hamid Ebadi, F Hernan Espinoza, Matt Fish, Qiang Gan, Benson M George, Jeffrey M Granja, Foad Green, Gunsagar S Gulati, Michael S Haney, Julie A Harris, Yanzhe He, Shayan Hosseinzadeh, Albin Huang, Kerwyn Casey Huang, Atsushi Iriki, Eric Jean, Kevin S Kao, Guruswamy Karnam, Aaron M Kershner, Bernhard M Kiss, William Kong, Maya E Kumar, Jonathan Lam, Song E Lee, Benoit Lehallier, Qiang Li, Yan Li, Ling Liu, Annie Lo, Wan-Jin Lu, Marisol F Lugo-Fagundo, Anjali Manjunath, Andrew P May, Ashley Maynard, Aaron McGeever, Madeleine McKay, Michael I Miller, Mais Moussa, Ravi Mylvaganam, EK Neumann, Joseph Noh, Roel Nusse, Irene Papatheodorou, Traci Peng, Lolita Penland, Katherine Pollard, Robert Puccinelli, Zhen Qi, Stephen R Quake, Thomas A Rando, Micha Sam Raredon, Karine Rizzoti, Katherine Rogers, Yanay Rosen, M Elizabeth Rothenberg, Meritxell Rovira, Yaroslava Ruzankina, Nicholas Schaum, Eran Segal, Jun Seita, Rahul Sinha, Rene V Sit, Justin Sonnenburg, Christof Stringer, Kai Tan, Michelle Tan, Sudhir Gopal Tattikota, Kyle J Travaglini, Carolina Tropini, Michelle Tsui, Lucas Waldburger, Bruce M Wang, Linda J van Weele, Brice M Weinstein, Michael N Wosczyna, Angela Wu, Jinyi Xiang, Sizun Xue, Kevin A Yamauchi, Andrew C Yang, Lakshmi P Yerra, Justin Youngyunpipatkul, Bo Yu, Fabio Zanini, Gizem Zardeneta, Tiffany Zee, Chunyu Zhao, Fan Zhang, Hui Zhang, Martin Jinye Zhang, Lu Zhou, and Daniel R Zollinger. The tabula sapiens: A multiple-organ, single-cell transcriptomic atlas of humans. *Science*, 376(6594):eabl4896, 2022. doi: 10.1126/science.abl4896. URL https://doi.org/10.1126/science.abl4896.

Christina V Theodoris, Ling Xiao, Anant Chopra, Mark D Chaffin, Zeina R Al Sayed, Matthew C Hill, Helene Mantineo, Elizabeth M Brydon, Zexian Zeng, X. Shirley Liu, and Patrick T Ellinor. Transfer learning enables predictions in network biology. *Nature*, 618(7965):616–624, 2023. doi: 10.1038/s41586-023-06139-9. URL https://doi.org/10.1038/s41586-023-06139-9.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 2017. URL https://proceedings.neurips.cc/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf.

Wenchuan Wang, Fan Yang, Yuan Fang, Duyu Tang, Junzhou Huang, Hui Lu, and Jianhua Yao. scbert: a large-scale pretrained deep language model for cell type annotation of single-cell rna-seq data. *bioRxiv*, 2021. doi: 10.1101/2021.12.05.471261. URL https://doi.org/10.1101/2021.12.05.471261. Preprint.

A. Wenteler, M. Occhetta, N. Branson, M. Huebner, V. Curean, W. T. Dee, W. T. Connell, A. Hawkins-Hooker, S. P. Chung, Y. Ektefaie, A. Gallagher-Syed, and C. M. V. Córdova. Pertevalscfm: Benchmarking single-cell foundation models for perturbation effect prediction. *bioRxiv*, 2024. doi: 10.1101/2024.10.02.616248. URL https://doi.org/10.1101/2024.10.02.616248. Preprint.

# A Additional Details

## A.1 Reproducibility

Training datasets are from publicly available sources (Pearce et al., 2025). Evaluation datasets from Tabula Sapiens v2 are available from the official repository (Tabula Sapiens Consortium et al., 2022).

Table 1: Hyperparameters used for model training.

| **Hyperparameter** | **Value** |
|---|---|
| *Model Architecture* | |
| Model dimension ($d_{\text{model}}$) | 512 |
| Number of transformer blocks ($n_{\text{blocks}}$) | 12 |
| Number of attention heads ($n_{\text{head}}$) | 8 |
| Feed-forward dimension ($d_{\text{hid}}$) | 1024 |
| Dropout | 0.1 |
| Context window | 1024 |
| *Training* | |
| Batch size | 512 |
| Gradient accumulation steps | 32 |
| Effective batch size | 16,384 |
| Epochs | 3 |
| Learning rate | $3 \times 10^{-4}$ |
| Weight decay | $10^{-5}$ |
| Optimizer betas | [0.9, 0.95] |
| Max gradient norm | 1.0 |
| Scheduler | Cosine annealing |
| Warmup steps | 2000 |
| Early stopping patience | 100 |
| *Data* | |
| Validation ratio | 0.001 |
| Test ratio | 0.001 |
| Masking probability ($p_{\text{mask}}$) | 0.5 |
| *Expression Encoding* | |
| Raw hidden dimension | 512 |
| Hard binning bins | 50 |
| Soft binning bins | 20 |
| Soft binning hidden dimension | 512 |
| Log binning max bins | 10 |
| *Other* | |
| Random seed | 777 |
| Number of data workers | 8 |

### A.2 HYPERPARAMETERS

Table 1 summarizes the hyperparameters used for training all models in this study. These hyperparameters were based on previous works (Cui et al., 2024; Adduri et al., 2025; Pearce et al., 2025).

### A.3 TRAINING DATASETS

Table 2 lists all 101 training datasets used in this study, including the dataset identifier and number of cells for each dataset. The total number of cells across all training datasets is 10,010,835.

### A.4 ADDITIONAL FIGURES

Table 2: Training datasets used for model pretraining. Dataset IDs are UUIDs from the curated dataset collection.

| Dataset ID | Cells | Dataset ID | Cells | Dataset ID | Cells |
|---|---|---|---|---|---|
| 7d98cc44-b090-4dc8-804f-2750c84fe9d7 | 2,489 | 2f05ab20-a092-4bab-9276-3e0eb24e3fee | 38,217 | 3966ba97-beb8-4d0b-9954-d3775cd2cd61 | 158,978 |
| 2d66790a-6621-4a49-8f0d-4002db5cc98d | 4,992 | 76150f40-1989-4977-9e23-696e72d59d9e | 118,672 | 28ab6eb8-dfa4-4536-9f26-7e06c1b98e8e | 25,741 |
| 50c4a6d6-940b-4c6a-a376-aea2ae2d3168 | 21,003 | c2a461b1-0c15-4047-9fcb-1f966fe55100 | 97,499 | c838aec3-03ef-4398-b882-0e3912abfff0 | 1,265,624 |
| cda2c8cd-be1c-42e5-b2cd-162caa1c4ce7 | 255,901 | e40c6272-af77-4a10-9385-62a398884f27 | 65,088 | 3a29c3df-b45a-403d-bd76-259640245432 | 4,992 |
| 19e46756-9100-4e01-8b0e-23b557558a4c | 66,985 | ed11cc3e-2947-407c-883c-c53b043917c3 | 8,573 | 61d327d1-2227-4c5f-9367-e3559dc79b07 | 796 |
| 70e4f35b-c98c-45a1-9aa9-2053b07315dd | 40,815 | 715327a6-7978-4896-ba91-69d6b04dbbfb | 40,191 | eeacb0c1-2217-4cf6-b8ce-1f0fedf1b569 | 9,337 |
| c5ac3ec2-24b0-43cc-9aab-bb0ebbe205ce | 4,992 | a6046b15-a095-43b0-9fb5-b36899d87fdb | 4,992 | 0fe5eed4-bcbc-4c00-a388-00bc1455a9b7 | 4,992 |
| bac09168-940b-4b11-8d55-b4955f80b98d | 12,461 | f9cfac8d-bff6-47a2-a1f6-503827d375f5 | 637 | 43aa19d2-c723-4822-979d-d2f0239835e0 | 37,121 |
| 095940cb-7422-4510-96e2-cbafd961eb88 | 52,045 | 30f5e171-83d7-4fc0-bf75-384f122346b3 | 1,790 | 3079e9b0-c6db-45c8-b998-a4555a73968e | 28,943 |
| 79884ae5-e026-4d4a-858f-e807960bd4f7 | 4,425 | 489318a0-24c3-4f5c-b105-f084ed0ea026 | 13,900 | 0920bcb8-4b3a-4e9d-a353-56f529fd3b32 | 48,478 |
| d551b400-b2e5-454d-b5a9-ece03e6b4739 | 4,992 | a7b4f565-691d-43ea-bf4a-d2d1d52bb4b4 | 27,111 | 93091496-be48-4122-b945-9af9c22a7535 | 28,718 |
| b9b592d4-a0cc-4694-8704-a6625829ef1f | 4,992 | 34229bdc-a895-4394-8820-574e4028d8c6 | 31,924 | 01209dce-3575-4bed-b1df-129f57fbc031 | 51,876 |
| 31f657dc-1875-4c4b-a5ca-ce63b3ef3a82 | 121,916 | e5f5d954-cf0e-4bd8-9346-8d1ddf15a08b | 2,487 | 879bb6df-cc2a-40f1-854b-5be9629d03b2 | 4,992 |
| edbfa04a-f1dd-496a-8237-df11d70621ca | 77,525 | 06ef6b36-6c9b-4e10-8a94-d0baf274276e | 10,533 | 79527108-1f6c-4152-afe0-1fcdf2e02ba3 | 4,992 |
| 6f0858c0-c590-4740-b022-c152e7608d66 | 4,992 | 9ea768a2-87ab-46b6-a73d-c4e915f25af3 | 40,268 | dee75ca4-8348-471e-bbeb-e2143209e3d0 | 4,992 |
| 949ec7fc-ac54-47c9-bb6f-ee9e67688cce | 38,937 | 0bae7ebf-eb54-46a6-be9a-3461cecefa4c | 27,675 | 1df5fa02-4a6e-4b00-a203-cb0a60e75637 | 4,992 |
| 729f397a-0812-4b52-a7d1-b377107ffb41 | 4,992 | 53d62b10-bae5-48ac-b16e-71be9ba6de59 | 4,992 | 364bd0c7-f7fd-48ed-99c1-ae26872b1042 | 931,012 |
| ae4552dc-e2ea-4d67-b375-03ec7480f780 | 37,275 | fe2479fd-daff-41a8-97dc-a50457ab1871 | 292,010 | 50eb1e23-b8d4-4f76-a184-44e5541fa05a | 4,775 |
| 77c1c785-809f-4065-8c54-6a0170783256 | 37,767 | 49108ba9-1b7a-4a8a-9859-3d32e6a83926 | 4,992 | 39ed7d98-676d-4b8d-9d0a-0f3b60914ead | 118,647 |
| dd03ce70-3243-4c96-9561-330cc461e4d7 | 23,732 | 00ba8341-48ec-4e4e-bb56-be0dd2dd7913 | 4,992 | 9ddea8d9-cc4c-420a-90f6-880996f808d4 | 100,307 |
| 9dbab10c-118d-496b-966a-67f1763a6b7d | 1,462,702 | ca421096-6240-4cee-8c12-d20899b3e005 | 81,736 | 731e6380-879f-4b0b-9a1f-2150208852ef | 2,065 |
| 43245158-5ae1-4e71-a9a6-67eef49c26bc | 113,304 | f801b7a9-80a6-4d09-9161-71474deb58ae | 6,044 | 54801477-ac3e-47e3-8170-96c5b40d5c10 | 4,992 |
| 43848156-ba94-47b5-8409-7535cea75678 | 4,992 | c5cfa2b7-abb1-4a50-908f-707b54ca606b | 14,094 | 1cf24082-59de-4029-ac81-6e398768af3a | 29,522 |
| 978a566a-dc27-478d-b306-26daba116c1f | 1,102,250 | 726afd49-df7b-4b56-967a-0fb79d85ee4b | 4,992 | 9adb1b29-65a2-4dd0-86bf-c02690d65fbd | 4,992 |
| e6361237-ac4e-4c5d-ad8f-f16aca0c0a8f | 66,719 | 344f27ab-428c-4a0e-a7e1-d4441f2f9b80 | 4,992 | af8b241a-c72c-4470-b1a4-80e7336c6ab6 | 4,335 |
| 346c5aad-b034-4248-8cbe-0a05fd634b9c | 163,779 | 529bb209-9d7b-44da-bfa4-f6e4745c46c2 | 32,678 | 33da10b0-9c1d-4c82-9b14-c67cdcf9fae5 | 30,022 |
| d98ea49e-b70d-4434-b850-bbe217c9b66e | 15,216 | 5e57cd50-8e42-42d6-940d-5c1660d06864 | 693,682 | 3e55180c-780a-4424-9434-5296640ffc0d | 7,774 |
| f3c49918-4707-4d92-bb6d-2b5b4eb9d1b4 | 15,177 | c5cddbbb-8ba4-4338-8b34-15edb5231e22 | 4,992 | b617ee1b-f8c8-4de9-b82b-e803ab93550d | 391,963 |
| 9f499d32-400d-4c42-ac9a-fb1481844fee | 56,367 | 93131426-0124-4ab4-a013-9dfbcd99d467 | 24,327 | abd889c6-f60a-4fbd-924e-ee1e9dcf909b | 4,992 |
| 2f6a20f1-173d-4b8d-860b-c47ffea120fa | 2,868 | 75548d10-160d-4f3e-b317-99ad9630c62d | 4,992 | e84f2780-51e8-4cfa-8aa0-13bbfef677c7 | 167,598 |
| be401db3-d732-408a-b0c4-71af0458b8ab | 135,462 | 6e9e3264-02e1-455a-840b-4fbcee132ae7 | 4,992 | 639ffc23-14db-4060-9027-3c90314200f8 | 35,290 |
| 4724c395-0c46-46d2-81f7-60fd271fb488 | 35,350 | 4fd2ee79-ab3a-4827-a773-1b7dcd099307 | 4,992 | c7775e88-49bf-4ba2-a03b-93f00447c958 | 647,366 |
| 15c5c186-df92-4b17-a253-199e10ffe98a | 4,992 | 019c7af2-c827-4454-9970-44d5e39ce068 | 12,590 | 8c42cfd0-0b0a-46d5-910c-fc833d83c45e | 65,662 |
| 1b767f95-d0a0-4a3d-b394-cc665d86c3dc | 34,933 | b3c55d0d-4529-4b61-b485-2902e6be0e4e | 4,992 | | |
| **Total: 10,010,835** | | | | | |

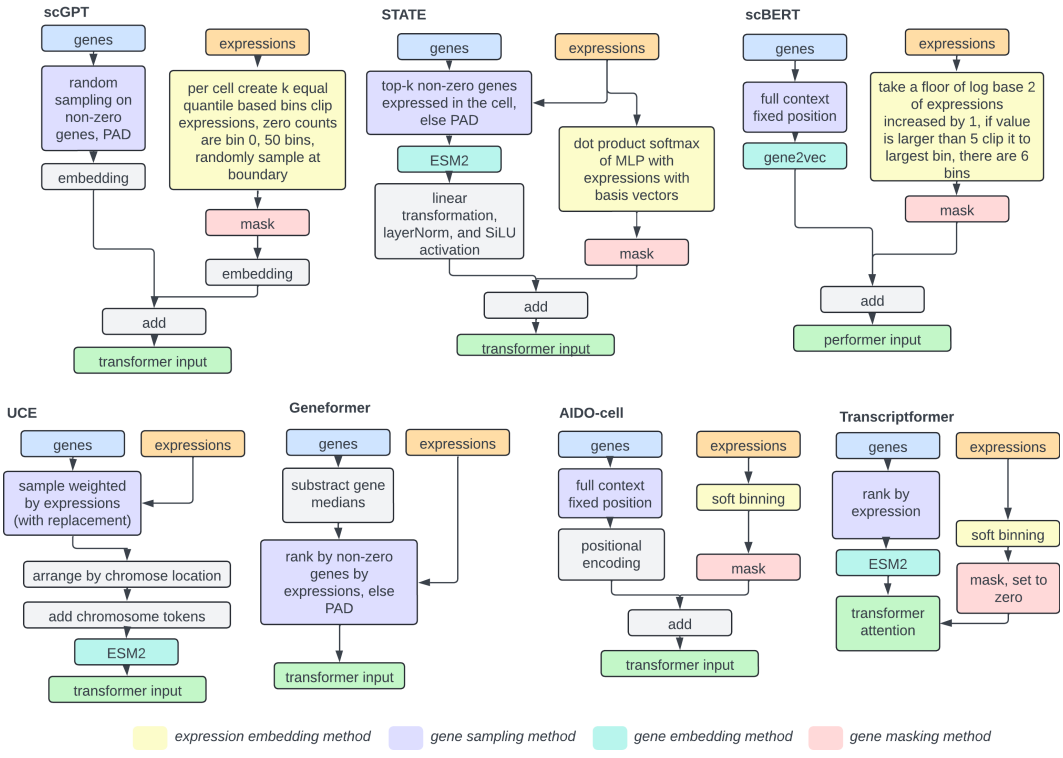


Figure 4: Schematic comparison of different gene and expression encoding methods across single-cell foundational models (Cui et al., 2024; Adduri et al., 2025; Ho et al., 2024; Wang et al., 2021; Theodoris et al., 2023; Pearce et al., 2025; Rosen et al., 2023).